

On moral responsibility

A metaphysical examination

Alexander Towell
lex@metafunctor.com

Contents

1	Introduction	1
2	Morality	2
3	Ascribing moral responsibility	12
4	Concluding thoughts	16

1 Introduction

People throughout history have believed that they belong to a special categorical class called *persons*. What supposedly makes members of this class special are their status as *moral agents*. But why do they warrant this distinction? Likely, it is based on criteria such as being able to make choices, form intentions, and perform actions. However, if the nature of reality is at odds with such criteria, our moral agency could be reduced to a pleasant fiction.

Metaphysics While challenges to the proposition that persons are moral agents have come from many different disciplines, few have posed a greater challenge to it than metaphysics. Metaphysics is a discipline in philosophy with a very long history. The etymology of *metaphysics* derives from the Greek words *metá* (after) and *physiká* (physics), and it simply refers to Aristotle's work that contiguously followed after his chapter on physics. However, over time the meaning of *metá* has been reinterpreted to mean something more substantial, like *beyond*. Thus, metaphysics is the study of that which is beyond physics. And arguably it is the most fundamental of all disciplines for it studies the categorical structure of reality itself.

There is some disagreement among meta-physicians about this categorical structure. Disagreements are not just over things like the kinds of categories that ought to be recognized, but they are also over things like whether reality even has a categorical structure. That is, do categories represent absolute ontological differences or do they represent the useful ways in which humans organize their experiences? Does reality have an inherent categorical structure (which may

someday be discovered) or do minds invent conventional categories to order their experiences?

Moral agency To say that we are moral agents is to say that we can be held morally responsible for our actions, which can run the gamut from *good* to *bad*. For example, people who are around a village during a volcanic eruption may demonstrate good moral conduct if they rescue villagers from the destructive power of the volcanic eruption. Alternatively, they may demonstrate bad moral conduct if they do nothing to help the villagers. The volcano eruption is causally responsible for the destruction to the village, but it is not morally responsible.

Under what circumstances are things morally responsible? A theory of moral responsibility might touch upon, at the very least, the following subjects:

1. The subject of morality itself.
2. The criteria for being a moral agent.
3. The circumstances under which moral responsibility can be ascribed.

In due course, after expanding upon each of these points, one might arrive at a satisfactory explanation that accounts for moral responsibility. However, the discipline of metaphysics has not only undermined the notion that persons are morally responsible, but also the notion that reality can even support the possibility of moral responsibility.

2 Morality

“Morality is simply the attitude we adopt towards people whom we personally dislike.” –Oscar Wilde

The etymology of *morality* derives from the Latin word *moralitas* (manner, character, proper behavior), and it can be used a) descriptively to refer to a code of conduct promoted by some group, or b) normatively to refer to a code of conduct that, given specified conditions, would be promoted by any moral agent. In what follows, our primary concern will be of the second normative type where the nature of morality can be explored in greater isolation from cultural norms.

What is the basis for morality? From a behavioral point of view, morality is arguably the study of right and wrong conduct. However, what is the basis of right and wrong? One possibility is that its basis is authoritative prescription, in which case morality is a dictum of *might makes right*—something is morally good or bad because an authority says so with no further explanation required. A forceful criticism of this type of moral theory is revealed by asking: is an action good because God commands it, or does God command it because it is good? If it is good simply because God commands it, then morality is nothing more than the arbitrary preference of God. But if God commands it because it is good, then morality has a basis independent of God.

Many would contend the basis is essentially rational; persons use reason and experience to inform their conduct, e.g., how can a group best realize common interests? In which case, a sense of right and wrong is, for the most part, the habituation of conduct that promotes cultural dominance or health, i.e., the principle of natural selection applied in the context of a social Darwinism.

While there are perhaps many other moral theories—some of which may not support the premise of moral responsibility—each must fall within one of two general categories: the conventional, and the intrinsic. On the one hand, if morality is conventional, then persons are only morally responsible by convention. On the other hand, if morality is categorically intrinsic, then arguably moral agents have an inherent duty to conform to it. But what does it mean to say that morality is categorically intrinsic?

Realism

There are many things—particulars—which are red. For instance, there are red cars, red apples, and red signs. But *what* makes them red? The philosophical discipline of realism submits that all red particulars are red in virtue of the existence of a Universal—a single abstract entity that is a part of all red particulars. More precisely, Universals are repeatable entities that can be exemplified simultaneously by different particulars, where entities are things like *properties*, like *red*, *relations*, like *behind*, and *kinds*, like *dog*.

Realists argue that attribute agreement is due to relationships between entities and Universals. Furthermore, they contend that the subject-predicate grammar of sentences reinforces this view, e.g., in the sentence “My apple is red,” the predicate “is red” seems to be picking out something (a noun)—call it *redness*—and linking it to the subject, “my apple.” Let us examine how a realist might account for following two sentences:

1. My apple is red.
2. My friend’s apple is red.

In each sentence, there are arguably two Universals being exemplified. First, there is the Universal property, *redness*, that is being exemplified by the two particulars, *my apple* and *my friend’s apple*. Second, there is the Universal kind, *apple-ness*, that is also being exemplified by *my apple* and *my friend’s apple*. The *redness* being exemplified by the entity, *my apple*, is numerically identical with the *redness* being exemplified by the entity, *my friend’s apple*. Likewise, the *apple-ness* that is being exemplified by *my apple* is numerically identical with the *appleness* that is being exemplified by *my friend’s apple*.

In the sentence, “This apple is red,” a metaphysical realist maintains that this denotes “*This apple* exemplifies *redness*,” where *redness* is a Universal that inheres in *this apple*. More generally, a realist claims that any subject-predicate sentence of the form “*a* is *F*” can be paraphrased as “*a* exemplifies *F-ness*” where *a* is an entity and *F-ness* is a Universal. In the same way, a realist could maintain that *morality*, like *color*, is a Universal. In the same way that an apple

can exemplify a color like *redness*, a person can exemplify a *virtue* (a moral excellence) like *honesty*, e.g., “Bob exemplifies honesty.” However, it turns out that the existence of Universals is one of the most important disputes in the discipline of metaphysics.

Paradox If “*a* is *F*” denotes “*a* exemplifies *F*-ness,” then we can construct a paradox. Let *F* mean “does not exemplify itself” and so *F*-ness could be the Universal property, *non-selfexemplification*. Thus, for some *a*, the proposition “*a* exemplifies *non-selfexemplification*,” should be either **true** or **false**. But the existence of this property leads immediately to a paradox, for if a thing does exemplify itself, then it does not exemplify itself and if a thing does not exemplify itself, then it does exemplify itself. It is self-contradictory.

On the one hand suppose that an entity, *a*, does exemplify itself. Then, since it is a property an entity exemplifies when it does not exemplify itself, *a* does not exemplify itself. On the other hand, suppose that *a* does not exemplify itself. Then, it does exemplify itself. It is similar to the sentence, “This sentence is a lie.” If it is **true**, then it is **false**; and if it is **false**, then it is **true**. It is self-contradictory.

To avoid this paradox, restrictions on what constitutes a Universal must be imposed, such as denying that *non-selfexemplification* is an actual Universal. However, as a result of this, the realist cannot take it for granted that “*a* is *F*” means that *F* must be selecting a Universal, *F*-ness. This is not necessarily a problem for morality, but the realist cannot simply claim, without qualification, that if a sentence can take on a subject-predicate form that necessarily means that the existence of Universals can be assumed to account for things like attribute-agreement. For instance, “*Socrates is honest*” does not necessarily mean that the adjective, *honest*, is pick out out a Universal entity, *honesty*.

A vicious regress? There is perhaps a deeper problem with the realist’s account. This problem can be illustrated by pointing out that in the statement, “*a* exemplifies *F*-ness,” the word *exemplifies* itself picks out a Universal—a repeatable entity that serves a *relational* role. After all, consider that two entities can enter into a numerically identical asymmetric relationship, *exemplification*. To explore this further, let us consider the following two sentences:

1. Socrates exemplifies courage.
2. Plato exemplifies honesty.

In each sentence above, there is the Universal entity, *exemplification*, that is being exemplified by the particulars, “Socrates exemplifies courage” and “Plato exemplifies honesty.” The *exemplification* exemplified by “Socrates exemplifies courage” is numerically identical with the *exemplification* exemplified by “Plato exemplifies honesty.”

More generally, if *exemplifies* picks out a Universal, then `generate_propositions(n)` as given by Algorithm 1 generates an arbitrary number of propositions for any sentence of the form “*a* is *F*.”

Example 1. Invoking `generate_propositions(3)` generates three proposi-

Input: An integer $n > 0$ indicating the number of propositions to generate

Output: A list of n propositions

```

 $p_1 \leftarrow$  “ $a$  exemplifies1  $F$ -ness”;
for  $i \leftarrow 2$  to  $n$  do
  |  $p_i \leftarrow$  “ $\{p_{i-1}\}$  exemplifies $i$  exemplification”;
end
return sequence of propositions  $p_1, p_2, \dots, p_n$ ;
Algorithm 1: Exemplification generator

```

tions, the third of which is $\{ a \text{ exemplifies}_1 F\text{-ness. } \} \text{ exemplifies}_2 \text{ exemplification. } \} \text{ exemplifies}_3 \text{ exemplification.}$

Moreover, `generate_propositions`(∞) generates an infinite regress of never-ending propositions and the worry is that this regress is vicious. Whether this is the case is certainly disputable, but one presumes there is a right answer and, if so, Realism’s veracity hinges on this issue.

The Bradley regress and its implications This problem, known as Bradley’s regress after F.H. Bradley who articulated a similar concern, strikes at the heart of the realist project. Bradley argued that relations cannot genuinely relate their relata without leading to infinite regress. If a relation R relates a to b , then we need to explain how R itself connects to a and b . This seems to require further relations R' and R'' , which in turn require further relations, and so on.

To make this concrete with moral properties: if Socrates exemplifies honesty, the realist says this involves three entities—Socrates (a particular), honesty (a Universal property), and exemplification (a Universal relation). But how does the relation of exemplification connect Socrates to honesty? It seems we need a further relation to connect exemplification to its relata, leading to an infinite hierarchy of relations.

Realist responses to the regress Realists have developed several strategies to address this challenge. Armstrong [1] argues that the regress is not vicious because exemplification is not a genuine Universal but rather a “non-relational tie” or “fundamental nexus” that requires no further explanation. On this view, that particulars instantiate Universals is a brute metaphysical fact—the basic furniture of reality includes not just objects and properties but also their instantiation.

An alternative response, defended by some contemporary realists, is to embrace the regress but deny its viciousness. They argue that while the regress generates an infinite series of truths, these truths are not increasingly fundamental. Rather, once we establish that a exemplifies F -ness, all the higher-order truths follow automatically without requiring separate metaphysical grounding. The truth of each proposition in the series is guaranteed by the truth of the first, making the regress benign rather than vicious.

A vicious infinite regress An infinite regress is vicious if each proposition p_n in the series requires the support of proposition p_{n+1} . If this is the case, then the infinite series of regressions required to provide such support can never be finished, e.g., p_1 requires p_2 , but p_2 requires p_3 , ad infinitum. Consequently, the realist's solution to problems like attribute agreement (" a is F " $\rightarrow$ " a exemplifies F -ness" where F -ness $\in$ set of Universal entities) would not be valid.

The regress becomes vicious specifically when:

1. Each level requires separate ontological grounding
2. The truth of level n depends on the prior establishment of level $n + 1$
3. No level is self-supporting or foundational

Implications for moral properties If the regress problem proves fatal to realism about Universals generally, it would be particularly damaging to moral realism. Moral properties like virtue, justice, and responsibility would lose their claimed objective status. The statement "This action is wrong" could not be grounded in the action's participation in the Universal property of wrongness, since the entire framework of Universals would be undermined.

This would push us toward either nominalism about moral properties (treating them as useful classifications without ontological import) or some form of moral anti-realism. The implications for moral responsibility are severe: if moral categories lack objective grounding, then the distinction between moral agents and non-agents may be merely conventional, a useful fiction we maintain for social purposes rather than a recognition of metaphysical fact.

Nominalism

There is competing philosophical doctrine to realism, called nominalism, which disclaims the existence of Universals. Nominalists believe that realism needlessly introduces the strange notion of Universals, exhibits incoherency, and is burdened by a busy ontology. In its place they propose a supposedly simpler but sufficient account for the apparent categorical structure of reality.

Nominalists submit that only actual particulars have independent existence. After all, how could a repeatable entity, like *redness*, simultaneously inhere in multiple entities, like an apple and a firetruck, such that the *redness* in the apple is numerically identical to the *redness* in the firetruck? The Realist is introducing a strange idea that entities with independent existence have non-local, unbounded physical presence. So, realists claim that the sentence, " a is F " denotes " a exemplifies F -ness" where F -ness is a Universal. In response, nominalists counter that " a is F " denotes " a is a member of the set of F things." For instance, "*Socrates* is courageous" denotes "*Socrates* is a member of the set of *courageous things*." Nominalists introduce *sets* (which are rigorously defined

mathematical entities) to do the work of Universals. And the primary pay-off to this is that they avoid having to invoke the strangeness of Universals.

Set An unordered collection of entities considered as a whole. The identity conditions for sets is such that, if a and b are sets, they are numerically identical iff for each entity in a that entity is in b . Using mathematical language, $a \equiv b$.

Abstract reference For simple subject-predicate sentences, the nominalist's account as already given would seem sufficient. However, how might they interpret a sentence including abstract parts, like "Honesty is a virtue?" The realist would contend that this can be understood to mean, "*Honesty* exemplifies being a *virtue*," *honesty* and *virtue* picking out abstract Universals. But what might nominalists say? One possible interpretation might be, "The set of *honest acts* is a subset of the set of *virtued acts*." While this gives nominalists a little trouble, their response might seem sufficient enough. However, their problems do not end here.

Set equivalency How might a nominalist deal with the supposition that all the things that have hearts also have kidneys and vice versa? In other words, the set of *things that have hearts* and set of *things that have kidneys* have the same members. But according to the identity conditions for sets, that means that the set of *things that have hearts* is identical to the set of *things that have kidneys*. Consequently, *having a heart* is the same thing as *having a kidney* (the set of *things that have hearts* $\equiv$ the set of *things that have kidneys*), which is plainly absurd.

What is the nominalist's response to this? Generally, they invoke the terminology of *possible* entities. Instead of " a is F " denoting " a is a member of the set of F -things," it denotes " a is a member of the set of *possible* F -things." The set of *possible things that have hearts* and set of *possible things that have kidneys* do not have the same members, therefore they are not equivalent. Unfortunately, part of the justification for nominalism was that it has a simpler ontology than realism, but the introduction of *possible worlds* somewhat undermines this effort.

Conventionalism Even if nominalism did offer a satisfactory account of attribute agreement and subject-predicate discourse, it is none-the-less nominative. Accordingly, there is no entity, *honesty*, there are only sets of resembling things. The nominalist posits that the sentence, "My friend is honest" simply means that *my friend* resembles other persons that are called honest. This does not appear to give ontological significance to *honesty*. Thus, we come back to where we started: if right and wrong is conventional, then is being a moral agent ontologically significant?

From metaphysical foundations to agency Having examined the metaphysical status of moral properties themselves, we must now turn to the agents who supposedly bear moral responsibility. Even if we were to resolve the debate between realists and nominalists about moral properties, a separate but equally fundamental question remains: what are the criteria for being a moral agent?

The metaphysical challenges we have explored regarding Universals and conventions now resurface in our examination of personhood and agency. If moral properties lack objective grounding, and if the agents themselves are metaphysically problematic, then the entire edifice of moral responsibility becomes questionable.

Criteria for moral agency

“I know only that what is moral is what you feel good after and what is immoral is what you feel bad after.” –Ernest Hemmingway

What is the criteria for being a moral agent? By definition, a moral agent is a being who can make moral judgments. Thus, at a minimum, it would seem a moral agent should have the capacity to make rational decisions—that is, reason about the effects their decisions have on reality.

Core requirements for agency The traditional philosophical account of moral agency requires several interconnected capacities. First, the agent must possess *rationality*—the ability to form beliefs based on evidence, draw logical inferences, and understand causal relationships. Second, the agent requires *intentionality*—the capacity to form intentions and direct actions toward specific goals. Third, the agent needs *reflective awareness*—the ability to consider one’s own mental states and evaluate one’s actions against moral principles.

Consider a concrete example: a physician deciding whether to administer an experimental treatment to a terminally ill patient. The physician must rationally evaluate the evidence for the treatment’s efficacy, form an intention based on the patient’s best interests, and reflect on the moral implications of the decision. Without any one of these capacities, genuine moral agency becomes questionable.

Degrees of moral responsibility Rather than treating moral agency as a binary property, we might consider it as existing along a spectrum. Young children, for instance, gradually develop moral agency as their cognitive capacities mature. They begin by understanding simple cause-and-effect relationships (“hitting hurts others”), progress to recognizing intentions (“I didn’t mean to break it”), and eventually develop the capacity for moral reflection (“Was my action fair?”).

This graduated view raises important questions about the boundaries of moral responsibility. If a person’s rational capacities are temporarily impaired—through intoxication, extreme emotional distress, or sleep deprivation—to what extent are they morally responsible for their actions? The law recognizes such gradations through concepts like diminished capacity and mitigating circumstances, but the philosophical foundations remain contentious.

The knowledge condition Beyond basic cognitive capacities, moral agency seems to require adequate knowledge of relevant circumstances. An agent cannot be held fully responsible for consequences they could not reasonably have foreseen. This epistemic requirement introduces complications: what constitutes “reasonable” foresight? Should moral agents be held responsible for their

ignorance itself?

Consider the historical example of physicians who prescribed mercury as medicine before its toxic effects were understood. Were they morally responsible for the resulting harm? The answer depends partly on whether the relevant knowledge was available and whether they had an obligation to seek it. This suggests that moral agency includes not just the capacity to make informed decisions but also the responsibility to become adequately informed.

So, how do persons measure up to these criteria? At the very least, they seem to have the capacity to reason about a limited domain of reality, and so are perhaps morally responsible within the limits circumscribed by that domain. But putting all of that aside, one can instead ask, “Do moral agents exist at all?” That is, are there subjects *making* decisions?

Persistence

To consider this question, let us analyze Chisholm’s essay, *Identity through Time* [2]. In this essay, he distinguishes two ways in which entities persist.

Intactly persisting entities In the first way, entities persist in a numerically identical sense such that an entity, a , at time t_1 is the same as entity, b , at time t_2 iff the set of parts for entity a at time t_1 is equivalent to the set of parts for entity b at time t_2 . Therefore, if $a \equiv b$, said entity persisted, at a minimum, $|t_2 - t_1|$ units of time.

Non-intactly persisting entities If an entity changes even a single part, it is not strictly the same. Therefore, since physical entities are in a constant state of flux, when one says physical entity a at time t_1 is the same as physical entity b at time t_2 (where $t_2 \neq t_1$), we do not mean that entities a and b are numerically identical. Rather, we mean to say that they are conventionally the same.

Ship of Theseus To clarify, let us consider the thought experiment of the *Ship of Theseus*, which raises the question of whether a ship which has had all its parts replaced retains the same identity. If the ship at t_1 has different parts than the ship at t_2 then they are not numerically identical; so, if the ships are identical, a non-intactly persisting sense of identity is being used. Chisholm contends that the question is not *is it the same ship* but *does it constitute the same ship?* And if it constitutes the same ship, that means it satisfies the Theory of Persistence.

Theory of Persistence A composite entity persists if conditions 1 and 2 are satisfied:

1. The precedent composite entity evolved from a antecedent composite entity. Note: Evolved denotes that after each successive change, the antecedent composite must have at least one part in common with the precedent composite.
2. The composite entity must satisfy additional quantifiable criteria.



In practice, how might this theory provide a solution for the *Ship of Theseus* dilemma? First, we must determine whether it satisfies, from initial state to final state, the *evolution* criterion. And second, it must satisfy additional criteria, e.g., *has the same sailing schedule*. For example, if the evolution of the ship undergoes fission then the *evolution* criterion by itself cannot determine which ship constitutes the *Ship of Theseus*. However, if only one of the ships after the fission has the same sailing schedule as the ship before the fission, then according to the *has the same sailing schedule* convention, that ship constitutes the *Ship of Theseus*.

Persons

It is not even clear how *persons* should be defined, but providing an unambiguous definition for it would help to resolve many controversial. For instance, to resolve the dispute over abortion, one could categorically determine when a bundle of cells constitutes person-hood. There appears to be a consensus that motile sperm cells and ova are not persons, but that newborn babies are. What was the critical turning point in which it from from just being a bundle of cells to being a person who has a bundle of cells? Expending considerable time and effort could be avoided if a definitive categorical account of this problem was given, but like with the pile of sand, such a definition seems problematically conventional.

Chisholm on persons In light of the *Theory of Persistence*, what can be said about persons? As with the *Ship of Theseus*, are person's identities also conventional? Or are they like sub-atomic particles which intactly persist? In *Which Physical Thing Am I* [3], Chisholm contends that *persons* are not conventional. In support of this, he asks you the reader to consider a hypothetical situation in which you are about to undergo an operation. The doctor presents you with two options: you can either choose the expensive option where you are subjected to total anesthesia during the operation, or you can choose the cheaper option. The cheaper option consists of the doctor giving you a pill before the operation to induce complete amnesia so that during the operation there is no memory of you, and giving you a pill after the operation so that everything that transpired during the operation is forgotten.

The question is, would it be reasonable for you to choose the cheap option? Chisholm contends that it would not be reasonable because you would believe that you would be the person enduring the operation. Furthermore, you would not be persuaded that this is not the case by any convention that denies that the person on the table is you. In other words, you would not consider your identity a matter of convention.

Critical response to Chisholm How might Chisholm's arguments be disputed? In the situation he presented, perhaps you would believe that you and the person enduring the operation in the cheaper option is the same person.

Furthermore, perhaps no convention could persuade you from this conviction. However, your conviction could simply be based on a convention that you deeply subscribe to. At this point, it does not seem like Chisholm has established that persons warrant special categorical distinction.

Bundle Theory and Ego Theory In the essay *Divided Minds and the Nature of Persons* [7], Parfit contrasts two competing views of persons, Bundle Theory and Ego Theory. Bundle Theory proposes that unity of consciousness cannot be explained by referring to persons; Bundle Theorists maintain that persons only exist as a feature of language. Persons are a convention—they are a *label* for a *bundle of things* like thoughts, memories, and sensations that are unified by causal kinds of relations. There are no subjects in which these bundles inhere. But the Ego Theorist proposes that these subjects really do exist—as subjects of experiences—and that they *have* a *bundle of things* like thoughts, memories, and sensations.

In support of Bundle Theory, Parfit presents the reader with a hypothetical situation in which a replica of you is created, atom by atom. At one end of the spectrum, if one-hundred percent of your atoms are replaced, this is clearly a *replica of you*. At the other end, if only one percent of your atoms are replaced, this is clearly *you*. Thus, in between these two extremes there must be a critical turning point where it goes from being *you* to being a *replica of you*. However, this appears implausible because any point chosen would seem arbitrary. How could a single atom make a difference? It is similar to the question, *When does sand become a pile of sand?* It is absurd to think that, at one point it is not a pile, and with the addition of a single grain of sand, it becomes a pile.

A practical illustration: The case of Phineas Gage Consider the famous case of Phineas Gage, a railroad construction foreman who survived an accident in 1848 where an iron rod passed through his skull, destroying much of his left frontal lobe. Before the accident, Gage was described as responsible, industrious, and well-liked. After recovering, he became impulsive, profane, and unable to hold steady employment. His friends said he was “no longer Gage.”

This case poses profound questions for both theories. For the Ego Theorist, is the post-accident person the same *subject of experience* as before? If Gage’s personality, values, and behavioral patterns changed so dramatically, what persisted that makes him the same person? The Ego Theorist must claim there is some essential “Gage-ness” that survived despite these changes.

The Bundle Theorist has a different challenge: if Gage is just a bundle of psychological states and dispositions, and these changed dramatically, in what sense is he the same person? The Bundle Theorist might respond that there is sufficient psychological continuity—memories, basic cognitive abilities, learned skills—to maintain identity, but the degree of change raises questions about where to draw the line.

Ego Theorists have difficulty responding to this question because they believe that persons are *subjects of experiences*, not just *bundles of things*. As such, there must be some special unknown thing that makes you who you are that may not be included in the replication. Bundle Theorists have little problem answering these questions, however, because they deny that there is a *you*—only

a *bundle of things*. If an exact copy of the *bundle of things* that constitutes you is made, what we choose to call this thing, *you* or a *replica of you*, is a matter of convention.

On split-brain studies Parfit believes that empirical results of split-brain case studies reinforce Bundle Theory. In this study, patients are shown a placard whose left half is blue and right half is red where each eye is allowed to see only one of the halves. Curiously, the indication is that patients end up expressing two separate streams of consciousness, one in which there is an awareness of seeing only blue and another in which there is an awareness of seeing only red. That seems to indicate that within each of these patients there are not one but two *subjects of experience*.

Problematically, Ego Theorists define a person as a *subject of experience* yet patients seem to be expressing two. Consequently, the Ego Theorist must now propose that *subjects of experiences* need not be persons at all, which seems like a difficult stretch to make. But in Bundle Theory, since persons are just *bundles of things* it does not seem odd that in this bundle there could be a pair of mental states, one in which there is an awareness of seeing only red, and another in which there is an awareness of seeing only blue. There is no subject in which awareness inheres.

The implications for moral responsibility These debates about personal identity and the nature of persons have profound implications for moral responsibility. If persons do not persist through time in any robust sense, or if they are merely conventional designations for bundles of experiences, then holding them responsible for past actions becomes problematic. How can we punish or praise a “person” for actions performed by what might be, metaphysically speaking, a different entity? The Ego Theorist’s subject of experience and the Bundle Theorist’s collection of mental states both face challenges in grounding the kind of persistent agency that moral responsibility seems to require.

3 Ascribing moral responsibility

“Responsibility: A detachable burden easily shifted to the shoulders of God, Fate, Fortune, Luck or ones neighbor. In the days of astrology it was customary to unload it upon a star.” –Ambrose Bierce, The Devil’s Dictionary, 1911

Under what circumstances can a moral agent’s conduct be ascribed moral responsibility? At a minimum, it would seem to require that its conduct be a product of being able to select, in a way that is unconstrained by causal agencies, a specific action among a set of possible actions—that is, a product of exercising a free will. For instance, an agent’s decision to go shopping could be a product of a free will if it was possible for the agent to have decided not to go shopping. Conversely, if an agent stumbles off a building and free falls towards the ground, the act of free falling is certainly not a product of free will. Under the circumstances, the agent had no choice but to fall.

Four-dimensionalism

It is often said that entities have location and extension in a three-dimensional space. That is, physical entities can be specified in a three-dimensional Cartesian coordinate system (see Figure 1), x , y , and z , which provides three spatial dimensions—respectively, length, height, and width, e.g., a cubic entity extends m units along the x , y , and z axes.




Figure 1: Three-dimensional Cartesian coordinate system with a unit cube illustrating spatial extension

However, our coordinate system is missing a crucial detail: at what time, t , will a particular entity be at a particular x , y , and z , and how long will it endure? Therefore, a physical entity's location and extension is more precisely specified in a four-dimensional coordinate system, x , y , z , and t , which provides three spatial dimensions and one temporal dimension, e.g., a hyper-cubic entity extends m units along the x , y , z , and t axes, or a person may extend an average of ≈ 0.6 meters along the x -axis, an average of ≈ 1.8 meters along the y -axis, an average of ≈ 0.3 meters along the z -axis, and ≈ 85.6 years along the t -axis.

Four-dimensional space-time has taken on a new-found importance with the advent of modern science. Under this model, when persons observe an entity at a particular time, they observe a single *time-slice* of it. More precisely, they observe a three-dimensional spatial “slice” (which is perpendicular to the temporal dimension) of a four-dimensional space-time entity. For example, when they observe a two-dimensional spatial slice of a sphere, they observe a disk; when they observe a three-dimensional spatial slice of a hyper-sphere, they observe a sphere.

At any time t , persons can only observe a three-dimensional spatial “slice” located at time t of a four-dimensional space-time entity. However, if they had access to a fifth-dimension, they could possibly rotate in such a way in five-dimensions to instantly observe the entirety of four-dimensional entities. From this “God’s eye” vantage point, the passage of time would represent a vector through a four-dimensional space-time.

Foreknowledge and choice Is this model of reality incompatible with moral agency? If free agents are immutable four-dimensional space-time entities, then would that not denote that their freedom is illusory? After all, they are unchanging—there does not appear to be an opportunity for choice, there is only the opportunity for existence. But perhaps this notion of timelessness is too challenging to our intuitions, so let us instead ask the question, does infallible foreknowledge remove the possibility for choice? In other words, given that (1) a free agent, given two options *a* and *b*, can freely choose between *a* and *b*, and (2) there is infallible foreknowledge pertaining to which option, *a* or *b*, the free agent will choose, are (1) and (2) compatible?

On the one hand, it seems infallible foreknowledge forces the free agent’s hand: the free agent must act in accordance with the foreknowledge. It has no choice; it cannot do otherwise. On the other hand, infallible knowledge about the future is arguably no more a limitation on choice than infallible knowledge about the past as long as that foreknowledge is not based on causal determination. And indeed, a four-dimensional space-time is not necessarily causally determined.

Causal Determinism

Supposing that reality is not an immutable four-dimensional space-time, let us turn our attention to the question of causal determination. Causal determinism is the doctrine that all events are the inevitable result of antecedent conditions, including the actions of *persons*. The totality of existence is a game of cosmic billiards where, given an initial state as input, the final state could be output by a computable function.

The Laplacian demon Pierre-Simon Laplace famously articulated the determinist position through his thought experiment of an omniscient intelligence—now known as Laplace’s demon. This hypothetical being, knowing the precise location and momentum of every atom in the universe at a given moment, could use the laws of physics to reconstruct the entire past and predict the entire future. In Laplace’s words, “nothing would be uncertain and the future, as the past, would be present to its eyes.”

This mechanistic view gained considerable support from classical physics, where Newton’s laws seemed to govern all motion with mathematical precision. Every effect had a cause, every cause had a prior cause, stretching back in an unbroken chain to the beginning of time. Human actions, being physical events in a physical brain, would be no exception to this universal causation.

Is causal determinism compatible with moral agency? In causal determinism, there is only one possible future; acts of apparent choice are the mechanical expression of a supposed moral agent’s heredity and past environment. Since its future is determined, that would denote that moral responsibility is not possible, for at the very least a moral agent must have the power to make choices between morally *good* and *bad* actions.

The incompatibilist position, defended by philosophers like van Inwagen [9], holds that determinism and free will cannot coexist. If our actions are the inevitable result of the laws of nature and events in the remote past—events over

which we had no control—then we cannot be morally responsible for them. This is formalized in van Inwagen’s Consequence Argument: no one has power over facts about the distant past; no one has power over the laws of nature; therefore, no one has power over facts about the future (including their own actions).

Compatibilist responses Not all philosophers accept this incompatibilist conclusion. Compatibilists argue that free will and moral responsibility are compatible with causal determinism, provided we understand these concepts correctly. Frankfurt [4] argues that what matters for moral responsibility is not the ability to do otherwise, but rather acting according to one’s own desires and values. In his famous thought experiments, Frankfurt presents cases where an agent cannot do otherwise (due to a neurological intervention that would trigger if they tried to choose differently) yet seems morally responsible because they act on their own reasons.

Consider a mundane example: when you choose coffee over tea this morning, a compatibilist would argue you acted freely if this choice flowed from your own preferences, beliefs, and deliberation—even if these were themselves determined by prior causes. The alternative—actions arising from no causes at all—would be random, not free. As Hume observed, liberty requires necessity; actions must flow from character to be attributable to an agent.

Quantum indeterminacy and libertarian free will The advent of quantum mechanics has introduced fundamental indeterminacy into our picture of reality. At the quantum level, events may be genuinely probabilistic rather than determined. Some libertarians about free will have seized on this as providing the metaphysical room needed for genuine choice. However, critics point out that random quantum events hardly seem a better foundation for moral responsibility than deterministic ones. If your decision to help or harm someone depends on quantum coin-flips in your neurons, how does this make you more responsible than if it depends on prior causes?

Robert Kane [5] has developed a sophisticated libertarian account that attempts to navigate between determinism and randomness. On his view, quantum indeterminacy may play a role in moments of moral conflict, where we experience genuine uncertainty about what to do. The resolution of this conflict, while not determined, is not merely random but shaped by our efforts of will. The indeterminacy provides the metaphysical space for self-formation, where we become responsible for our character through these undetermined but non-random choices.

An argument from absurdity: the coherency of choice For the sake of argument, assume choice is not an illusion. How then might we offer a rational account of it? But if something is rational, does that not signify that it can be explicated in terms of its antecedent conditions? How else can a physical process be rationally explained if not by cause-and-effect? But if choice can be understood in terms of its antecedent conditions, then it is not free—it was determined by cause-and-effect.

This dilemma—that free will seems impossible whether determinism is true or false—has led some philosophers to hard incompatibilism. Derk Pereboom and others argue that free will is incompatible both with determinism and with

indeterminism, and therefore moral responsibility in the traditional sense is impossible. Yet even these skeptics often argue that we can maintain many moral practices and attitudes without ultimate moral responsibility. We can still shape behavior through praise and blame, protect society from dangerous individuals, and express moral emotions like guilt and indignation—understanding these as forward-looking practices rather than backward-looking judgments of ultimate desert.

4 Concluding thoughts

The metaphysical challenges to moral responsibility examined in this paper reveal deep tensions in our understanding of persons as moral agents. We have seen how the very foundations of moral agency—the existence of persons as persistent subjects, the reality of choice, and the possibility of acting otherwise—face serious philosophical challenges.

The cumulative challenge Each metaphysical issue compounds the others. If Universals do not exist and morality is merely conventional, then moral categories lack ontological grounding. If persons are merely bundles of experiences with no persistent identity, then who exactly bears moral responsibility across time? If our sense of temporal passage is illusory and we exist as static four-dimensional worms, or if causal determinism governs all events, then the notion of choice—essential to moral agency—becomes problematic.

The split-brain studies discussed by Parfit raise particularly troubling questions. If a single body can house multiple streams of consciousness, and if personal identity is conventional rather than categorical, then the assignment of moral responsibility becomes a matter of pragmatic decision rather than metaphysical fact. We might hold “persons” responsible not because they are ultimate sources of action, but because doing so serves useful social functions.

Pragmatic implications Yet the practical question remains: what would be entailed by widespread acceptance that moral agency is not a legitimate concept? Historical evidence suggests that societies do not simply collapse when confronted with challenges to their fundamental assumptions. The scientific revolution removed Earth from the center of the universe without ending human civilization. Darwin’s theory of evolution challenged human exceptionalism without eliminating ethics.

Similarly, the metaphysical challenges to moral responsibility need not lead to nihilism or social breakdown. Even hard incompatibilists like Pereboom argue for maintaining many moral practices on pragmatic grounds. We might reconceive punishment as rehabilitation and deterrence rather than retribution. We might understand moral emotions as evolutionary adaptations that promote cooperation rather than responses to ultimate desert. We might treat character development as a causal process subject to influence rather than a product of libertarian free will.

The persistence of moral phenomenology Perhaps most significantly, the phenomenology of moral agency—the felt experience of deliberation, choice,

and responsibility—persists regardless of metaphysical arguments. Even those intellectually convinced that free will is an illusion continue to experience decision-making as involving genuine alternatives. This experiential reality may be more fundamental to moral practice than metaphysical truth.

P.F. Strawson [8] argued that our practices of holding people responsible are grounded not in metaphysical facts but in reactive attitudes like resentment, gratitude, and indignation. These attitudes are constitutive of moral responsibility rather than responses to it. On this view, the question is not whether people *are* moral agents in some ultimate metaphysical sense, but whether treating them as such is integral to the forms of life we value.

Remaining questions This investigation leaves several questions unresolved. If moral responsibility is a useful fiction, is there an ethical obligation to maintain or dispel this fiction? Can a society function with full awareness that its moral categories are conventional constructions? Does the truth value of moral agency even matter if its practical effects remain constant?

Furthermore, the metaphysical issues themselves remain contested. Perhaps future developments in physics, neuroscience, or philosophy will resolve the puzzles of personal identity, free will, and moral categories. Or perhaps these questions, like Kant’s antinomies, reveal the limits of human reason when applied to fundamental aspects of existence.

Final reflection The metaphysical examination of moral responsibility ultimately returns us to a fundamental question about the relationship between truth and practice. Even if persons are conventional designations, even if the universe is deterministic, even if choice is illusory, we continue to live as if we are moral agents. This may be because moral agency, like consciousness itself, is not something we *have* but something we *are*—a mode of being that precedes and shapes our theoretical understanding rather than following from it.

Whether this constitutes a noble lie, a pragmatic necessity, or a deeper truth about the nature of human existence remains an open question. What seems clear is that the concept of moral responsibility, whatever its ultimate metaphysical status, remains central to human self-understanding and social organization. The challenge is not to abandon this concept in the face of metaphysical difficulties, but to refine our understanding of it in light of these challenges.

References

- [1] David M Armstrong. *Universals: An opinionated introduction*. Westview Press, 1989.
- [2] Roderick M Chisholm. Identity through time. In Howard E Kiefer and Milton K Munitz, editors, *Language, Belief, and Metaphysics*, pages 163–182. State University of New York Press, 1970.
- [3] Roderick M Chisholm. Which physical thing am i? an excerpt from “is there a mind-body problem?”. *Metaphysics: The Big Questions*, pages 296–300, 1990.

- [4] Harry G Frankfurt. Alternate possibilities and moral responsibility. *The Journal of Philosophy*, 66(23):829–839, 1969.
- [5] Robert Kane. *The significance of free will*. Oxford University Press, 1996.
- [6] Michael J Loux. *Metaphysics: A contemporary introduction*. Routledge, 3 edition, 2006.
- [7] Derek Parfit. Divided minds and the nature of persons. *Mindwaves*, pages 19–28, 1987.
- [8] Peter F Strawson. *Freedom and resentment and other essays*. Methuen, 1974.
- [9] Peter van Inwagen. The incompatibility of free will and determinism. *Philosophical Studies*, 27(3):185–199, 1975.